

Kafka Producer Advanced

Objectives Create Producer

- ❖ Cover advanced topics regarding Java Kafka Producer
- ❖ Custom Serializers
- ❖ Custom Partitioners
- ❖ Batching
- ❖ Compression
- ❖ Retries and Timeouts
- ❖ Message serialization using Avro

Kafka Producer

Tos

- ❖ Kafka client that publishes records to Kafka cluster
- ❖ Thread safe
- ❖ Producer has pool of buffer that holds to-be-sent records
 - ❖ background I/O threads turning records into request bytes and transmit requests to Kafka
- ❖ Close producer so producer will not leak resources

Kafka Producer Send, Acknowledgments and Buffers

- ❖ *send()* method is asynchronous
 - ❖ adds the record to output buffer and return right away
 - ❖ buffer used to batch records for efficiency IO and compression
- ❖ `acks` config controls Producer record durability. "all" setting ensures full commit of record, and is most durable and least fast setting
- ❖ Producer can retry failed requests
- ❖ Producer has buffers of unsent records per topic partition (sized at ***batch.size***)

Kafka Producer: Buffering and batching

- ❖ Kafka Producer buffers are available to send immediately as fast as broker can keep up (limited by inflight ***max.in.flight.requests.per.connection***)
- ❖ To reduce requests count, set ***linger.ms*** > 0
 - ❖ wait up to ***linger.ms*** before sending or until batch fills up whichever comes first
 - ❖ Under heavy load ***linger.ms*** not met, under light producer load used to increase broker IO throughput and increase compression
- ❖ ***buffer.memory*** controls total memory available to producer for buffering
 - ❖ If records sent faster than they can be transmitted to Kafka then this buffer gets exceeded then additional send calls block. If period blocks (***max.block.ms***) after then Producer throws a **TimeoutException**

Producer Ack

- ❖ Producer Config property ***acks***
 - ❖ **(default all)**
- ❖ Write Acknowledgment received count required from partition leader before write request deemed complete
- ❖ Controls **Producer** sent records durability
- ❖ Can be all (-1), none (0), or leader (1)

Acks 0 (NONE)

- ❖ acks=0
- ❖ Producer does not wait for any ack from broker at all
- ❖ Records added to the socket buffer are considered sent
- ❖ No guarantees of durability - maybe
- ❖ Record Offset returned is set to -1 (unknown)
- ❖ Record loss if leader is down
- ❖ Use Case: maybe log aggregation

Acks 1 (LEADER)

- ❖ acks=1
- ❖ Partition leader wrote record to its local log but responds without followers confirmed writes
- ❖ If leader fails right after sending ack, record could be lost
 - ❖ Followers might have not replicated the record
- ❖ Record loss is rare but possible
- ❖ Use Case: log aggregation

Acks -1 (ALL)

- ❖ acks=all or acks=-1
- ❖ Leader gets write confirmation from full set of ISRs before sending ack to producer
- ❖ Guarantees record not be lost as long as one ISR remains alive
- ❖ Strongest available guarantee
- ❖ Even stronger with broker setting ***min.insync.replicas*** (specifies the minimum number of ISRs that must acknowledge a write)
- ❖ Most Use Cases will use this and set a ***min.insync.replicas > 1***

Kafka Producer config Acks

```
c StockPriceKafkaProducer.java x
StockPriceKafkaProducer
17 > public class StockPriceKafkaProducer {
18     private static final Logger logger = LoggerFactory.getLogger(Stock
19
20
21     private static Producer<String, StockPrice> createProducer() {
22         final Properties props = new Properties();
23         setupBootstrapAndSerializers(props);
24         setupBatchingAndCompression(props);
25         setupRetriesInFlightTimeout(props);
26
27
28         //Set number of acknowledgments - acks - default is all
29         props.put(ProducerConfig.ACKS_CONFIG, "all");
30
31         return new KafkaProducer<>(props);
32     }
}
```

Producer Buffer Memory Size

- ❖ Producer config property: ***buffer.memory***
 - ❖ ***default 32MB***
- ❖ Total memory (bytes) producer can use to buffer records to be sent to broker
- ❖ Producer blocks up to ***max.block.ms*** if ***buffer.memory*** is exceeded
 - ❖ if it is sending faster than the broker can receive, exception is thrown

Batching by Size

- ❖ Producer config property: ***batch.size in*** total bytes
 - ❖ Default 16K
- ❖ Producer batch records
 - ❖ fewer requests for multiple records sent to same partition
 - ❖ Improved IO throughput and performance on both producer and server
- ❖ If record is larger than the batch size, it will not be batched
- ❖ Producer sends requests containing multiple batches
 - ❖ batch per partition
- ❖ Small batch size reduce throughput and performance. If batch size is too big, memory allocated for batch is wasted

Batching by Time and Size - 1

- ❖ *Producer* config property: *linger.ms*
 - ❖ Default 0
- ❖ Producer groups together any records that arrive before they can be sent into a batch
 - ❖ good if records arrive faster than they can be sent out
- ❖ *Producer* can reduce requests count even under moderate load using *linger.ms*

Batching by Time and Size - 2

- ❖ *linger.ms* adds delay to wait for more records to build up so larger batches are sent
 - ❖ *good brokers throughput at cost of producer latency*
- ❖ If **producer** gets records whose size is **batch.size** or more for a broker's leader partitions, then it is sent right away
- ❖ If **Producers** gets less than **batch.size** but *linger.ms* interval has passed, then records for that partition are sent
- ❖ Increase to improve throughput of Brokers and reduce broker load (common improvement)

Compressing Batches

- ❖ *Producer* config property: ***compression.type***
 - ❖ Default 0
 - ❖ Producer compresses request data
 - ❖ By default producer does not compress
 - ❖ Can be set to none, gzip, snappy, or lz4
 - ❖ Compression is by batch
 - ❖ improves with larger batch sizes
- ❖ End to end compression possible if Broker config “*compression.type*” set to producer. Compressed data from producer sent to log and consumer by broker

Batching and Compression

```
14 > public class StockPriceKafkaProducer {  
15  
16     private static Producer<String, StockPrice> createProducer() {  
17         final Properties props = new Properties();  
18         setupBootstrapAndSerializers(props);  
19         setupBatchingAndCompression(props);
```

```
57  
58     private static void setupBatchingAndCompression(Properties props) {  
59         //Wait up to 50 ms to batch to Kafka - linger.ms  
60         props.put(ProducerConfig.LINGER_MS_CONFIG, 200);  
61  
62         //Holds up to 64K per partition default is 16K - batch.size  
63         props.put(ProducerConfig.BATCH_SIZE_CONFIG, 16_384 * 4);  
64  
65         //Holds up to 64 MB default is 32MB for all partition buffers  
66         // - "buffer.memory"  
67         props.put(ProducerConfig.BUFFER_MEMORY_CONFIG, 33_554_432 * 2);  
68  
69         //Set compression type to snappy - compression.type  
70         props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "snappy");  
71     }
```

Custom Serializers

- ❖ You don't have to use built in serializers
- ❖ You can write your own
- ❖ Just need to be able to convert to/fro a byte[]
- ❖ Serializers work for keys and values
- ❖ *value.serializer* and *key.serializer*

Custom Serializers Config

```
14 ► public class StockPriceKafkaProducer {  
15  
16     private static Producer<String, StockPrice> createProducer() {  
17         final Properties props = new Properties();  
18         setupBootstrapAndSerializers(props);  
19  
30  
31     private static void setupBootstrapAndSerializers(Properties props) {  
32         props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,  
33             StockAppConstants.BOOTSTRAP_SERVERS);  
34         props.put(ProducerConfig.CLIENT_ID_CONFIG, "StockPriceKafkaProducer");  
35         props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,  
36             StringSerializer.class.getName());  
37  
38         //Custom Serializer - config "value.serializer"  
39         props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,  
40             StockPriceSerializer.class.getName());  
41     }
```

Custom Serializer

Tos

```
StockPriceSerializer.java x

1 package com.cloudurable.kafka.producer;
2
3 import com.cloudurable.kafka.producer.model.StockPrice;
4 import org.apache.kafka.common.serialization.Serializer;
5
6 import java.nio.charset.StandardCharsets;
7 import java.util.Map;
8
9 public class StockPriceSerializer implements Serializer<StockPrice> {
10
11     @Override
12     public byte[] serialize(String topic, StockPrice data) {
13         return data.toJson().getBytes(StandardCharsets.UTF_8);
14     }
15
16     @Override
17     public void configure(Map<String, ?> configs, boolean isKey) {
18     }
19
20     @Override
21     public void close() {
22     }
23 }
24
```

Stock Price

Tos

```
c StockPrice.java x
StockPrice
1 package com.cloudurable.kafka.producer.model;
2
3 import io.advantageous.boon.json.JsonFactory;
4
5 public class StockPrice {
6
7     private final int dollars;
8     private final int cents;
9     private final String name;
10
11    public String toJson() {
12        return "{" +
13            "\"dollars\": " + dollars +
14            ", \"cents\": " + cents +
15            ", \"name\": \"\" + name + '\"' +
16            "}";
17    }
18}
```

Broker Follower Write Timeout

- ❖ **Producer** config property: *request.timeout.ms*
 - ❖ Default 30 seconds (30,000 ms)
- ❖ Maximum time broker waits for confirmation from followers to meet Producer acknowledgment requirements for **ack=all**
- ❖ Measure of broker to broker latency of request
- ❖ 30 seconds is high, long process time is indicative of problems

Producer Retries

- ❖ Producer config property: ***retries***
 - ❖ ***Default 0***
- ❖ Retry count if ***Producer*** does not get ack from Broker
 - ❖ only if record send fail deemed a transient error ([API](#))
 - ❖ as if your producer code resent record on failed attempt
- ❖ timeouts are retried, ***retry.backoff.ms*** (default to 100 ms) to wait after failure before ***retry***

Retry, Timeout, Back-off Example

```
13
14 > public class StockPriceKafkaProducer {
15
16     private static Producer<String, StockPrice> createProducer() {
17         final Properties props = new Properties();
18         setupBootstrapAndSerializers(props);
19         setupBatchingAndCompression(props);
20         setupRetriesInFlightTimeout(props);

40
41     private static void setupRetriesInFlightTimeout(Properties props) {
42
43         //Only two in-flight messages per Kafka broker connection
44         // - max.in.flight.requests.per.connection (default 5)
45         props.put(ProducerConfig.MAX_IN_FLIGHT_REQUESTS_PER_CONNECTION,
46                 1);
47
48         //Set the number of retries - retries
49         props.put(ProducerConfig.RETRIES_CONFIG, 2);
50
51         //Request timeout - request.timeout.ms
52         props.put(ProducerConfig.REQUEST_TIMEOUT_MS_CONFIG, 15_000);
53
54         //Only retry after one second.
55         props.put(ProducerConfig.RETRY_BACKOFF_MS_CONFIG, 1_000);
56
57 }
```

Producer Partitioning

- ❖ **Producer** config property: ***partitioner.class***
 - ❖ org.apache.kafka.clients.producer.internals.DefaultPartitioner
- ❖ Partitioner class implements [Partitioner](#) interface
- ❖ Default Partitioner partitions using hash of key if record has key
- ❖ Default Partitioner partitions uses round-robin if record has no key

Configuring Partitioner

```
c StockPriceKafkaProducer.java x
StockPriceKafkaProducer createProducer()
17 > public class StockPriceKafkaProducer {
18     private static final Logger logger = LoggerFactory.getLogger(Sto
19
20
21     private static Producer<String, StockPrice> createProducer() {
22         final Properties props = new Properties();
23         setupBootstrapAndSerializers(props);
24         setupBatchingAndCompression(props);
25         setupRetriesInFlightTimeout(props);
26
27         //Install partitioner -- "partitioner.class"
28         props.put(ProducerConfig.PARTITIONER_CLASS_CONFIG,
29                   StockPricePartitioner.class.getName());
30
31         props.put("importantStocks", "IBM,UBER");
```

StockPricePartitioner

```
StockPriceKafkaProducer.java x StockPricePartitioner.java x
1 package com.cloudurable.kafka.producer;
2
3 import org.apache.kafka.clients.producer.Partitioner;
4 import org.apache.kafka.common.Cluster;
5 import org.apache.kafka.common.PartitionInfo;
6
7 import java.util.*;
8
9 public class StockPricePartitioner implements Partitioner{
10
11     private final Set<String> importantStocks;
12     public StockPricePartitioner() { importantStocks = new HashSet<>(); }
13
14     @Override
15     public int partition(final String topic,
16                         final Object objectKey,
17                         byte[] keyBytes, final Object value,
18                         final byte[] valueBytes,
19                         final Cluster cluster) {...}
20
21     @Override
22     public void close() {
23     }
24
25     @Override
26     public void configure(Map<String, ?> configs) {...}
27 }
```

StockPricePartitioner Partition()

```
StockPriceKafkaProducer.java x StockPricePartitioner.java x
StockPricePartitioner partition()
17     public int partition(final String topic,
18                           final Object objectKey,
19                           final byte[] keyBytes,
20                           final Object value,
21                           final byte[] valueBytes,
22                           final Cluster cluster) {
23
24     final List<PartitionInfo> partitionInfoList =
25         cluster.availablePartitionsForTopic(topic);
26     final int partitionCount = partitionInfoList.size();
27     final int importantPartition = partitionCount -1;
28     final int normalPartitionCount = partitionCount -1;
29
30     final String key = ((String) objectKey);
31
32     if (importantStocks.contains(key)) {
33         return importantPartition;
34     } else {
35         return Math.abs(key.hashCode()) % normalPartitionCount;
36     }
37
38 }
```

StockPricePartitioner

```
StockPriceKafkaProducer.java × StockPricePartitioner.java ×
StockPricePartitioner configure()
39
40     @Override
41     public void close() {
42     }
43
44     @Override
45     public void configure(Map<String, ?> configs) {
46         final String importantStocksStr = (String) configs.get("importantStocks");
47         Arrays.stream(importantStocksStr.split(regex: ","))
48             .forEach(importantStocks::add);
49     }
50 }
```

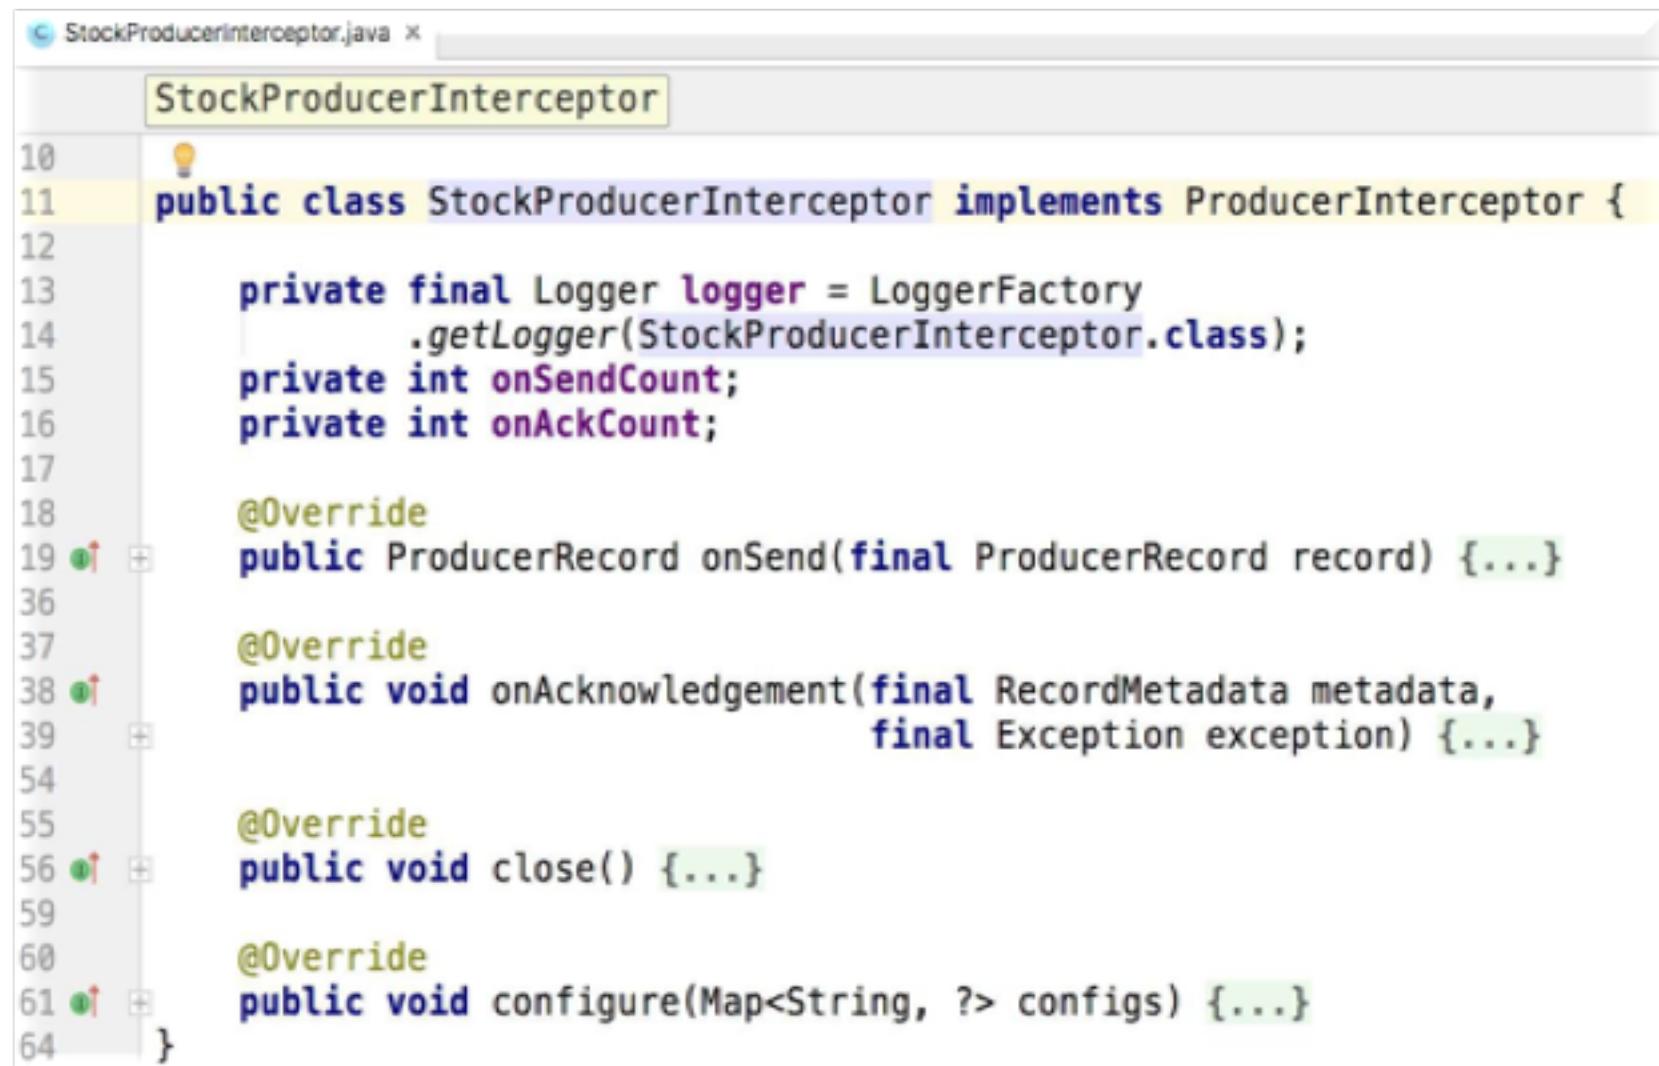
Producer Interception

- ❖ **Producer** config property: *interceptor.classes*
 - ❖ empty (you can pass an comma delimited list)
- ❖ interceptors implementing [ProducerInterceptor](#) interface
- ❖ intercept records producer sent to broker and after acks
- ❖ you could mutate records

Kafka Producer - Interceptor

```
StockPriceKafkaProducer.java x
StockPriceKafkaProducer getStockSenderList()
17 > public class StockPriceKafkaProducer {
18     private static final Logger logger = LoggerFactory.getLogger(StockPriceKafkaProducer.class);
19
20
21     private static Producer<String, StockPrice> createProducer() {
22         final Properties props = new Properties();
23         setupBootstrapAndSerializers(props);
24         setupBatchingAndCompression(props);
25         setupRetriesInFlightTimeout(props);
26
27         //Install interceptor list - config "interceptor.classes"
28         props.put(ProducerConfig.INTERCEPTOR_CLASSES_CONFIG,
29                   StockProducerInterceptor.class.getName());
30
31         //Set number of acknowledgments - acks - default is all
32         props.put(ProducerConfig.ACKS_CONFIG, "all");
33
34         return new KafkaProducer<>(props);
35     }
}
```

KafkaProducer ProducerInterceptor



```
StockProducerInterceptor.java x
StockProducerInterceptor
10
11 public class StockProducerInterceptor implements ProducerInterceptor {
12
13     private final Logger logger = LoggerFactory
14         .getLogger(StockProducerInterceptor.class);
15     private int onSendCount;
16     private int onAckCount;
17
18     @Override
19     public ProducerRecord onSend(final ProducerRecord record) {...}
36
37     @Override
38     public void onAcknowledgement(final RecordMetadata metadata,
39                                  final Exception exception) {...}
54
55     @Override
56     public void close() {...}
59
60     @Override
61     public void configure(Map<String, ?> configs) {...}
64 }
```

ProducerInterceptor onSend



```
StockProducerInterceptor.java ×
18
19 ❶ StockProducerInterceptor
20
21     @Override
22     public ProducerRecord onSend(final ProducerRecord record) {
23         onSendCount++;
24         if (logger.isDebugEnabled()) {
25             logger.debug(String.format("onSend topic=%s key=%s value=%s %d \n",
26                                     record.topic(), record.key(), record.value().toString(),
27                                     record.partition()
28                                     ));
29         } else {
30             if (onSendCount % 100 == 0) {
31                 logger.info(String.format("onSend topic=%s key=%s value=%s %d \n",
32                                     record.topic(), record.key(), record.value().toString(),
33                                     record.partition()
34                                     ));
35         }
36     }
37     return record;
38 }
```

Output

ic=stock-prices2 key=UBER value=StockPrice{dollars=737, cents=78, name='

ProducerInterceptor onAck

```
c StockProducerInterceptor.java x
StockProducerInterceptor onAcknowledgement()
37     @Override
38     public void onAcknowledgement(final RecordMetadata metadata,
39                                   final Exception exception) {
40         onAckCount++;
41
42         if (logger.isDebugEnabled()) {
43             logger.debug(String.format("onAck topic=%s, part=%d, offset=%d\n",
44                                     metadata.topic(), metadata.partition(), metadata.offset())
45         );
46     } else {
47         if (onAckCount % 100 == 0) {
48             logger.info(String.format("onAck topic=%s, part=%d, offset=%d\n",
49                                     metadata.topic(), metadata.partition(), metadata.offset())
50         );
51     }
52 }
53 }
```

ProducerInterceptor the rest

```
c StockProducerInterceptor.java x
StockProducerInterceptor configure()
48     logger.info(String.format("onAck topic=%s",
49                     metadata.topic(), metadata.partition
50             ));
51     }
52   }
53 }
54
55 @Override
56 public void close() {
57 }
58
59 @Override
60 public void configure(Map<String, ?> configs) {
61 }
62 }
```

KafkaProducer send() Method

- ❖ Two forms of send with callback and with no callback both return Future
 - ❖ Asynchronously sends a record to a topic
 - ❖ Callback gets invoked when send has been acknowledged.
- ❖ send is asynchronous and return right away as soon as record has added to send buffer
- ❖ Sending many records at once without blocking for response from Kafka broker
- ❖ Result of send is a RecordMetadata
 - ❖ record partition, record offset, record timestamp
- ❖ Callbacks for records sent to same partition are executed in order

KafkaProducer send() Exceptions

- ❖ ***InterruptedException*** - If the thread is interrupted while blocked ([API](#))
- ❖ ***SerializationException*** - If key or value are not valid objects given configured serializers([API](#))
- ❖ ***TimeoutException*** - If time taken for fetching metadata or allocating memory exceeds max.block.ms, or getting acks from Broker exceed timeout.ms, etc. ([API](#))
- ❖ ***KafkaException*** - If Kafka error occurs not in public API. ([API](#))

Using send method

```
c StockSender.java x
StockSender run()
51
52
53
54
55
56    try {
        final Future<RecordMetadata> future = producer.send(record);
        if (sentCount % 100 == 0) {
            displayRecordMetaData(record, future);
    }
```

```
c StockSender.java x
StockSender displayRecordMetaData()
74    private void displayRecordMetaData(final ProducerRecord<String, StockPrice> record,
75                                         final Future<RecordMetadata> future)
76                                         throws InterruptedException, ExecutionException {
77     final RecordMetadata recordMetadata = future.get();
78     logger.info(String.format("\n\t\tkey=%s, value=%s " +
79                           "\n\t\tsent to topic=%s part=%d off=%d at time=%s",
80                           record.key(),
81                           record.value().toJson(),
82                           recordMetadata.topic(),
83                           recordMetadata.partition(),
84                           recordMetadata.offset(),
85                           new Date(recordMetadata.timestamp())
86                           ));
87 }
```

KafkaProducer flush() method

flush() method sends all buffered records now (even if *linger.ms* > 0)

- ❖ blocks until requests complete
- ❖ Useful when consuming from some input system and pushing data into Kafka
- ❖ **flush()** ensures all previously sent messages have been sent
 - ❖ you could mark progress as such at completion of flush

KafkaProducer close()

- ❖ close() closes producer
 - ❖ frees resources (threads and buffers) associated with producer
- ❖ Two forms of method
- ❖ both block until all previously sent requests complete or duration passed in as args is exceeded
- ❖ close with no params equivalent to close(Long.MAX_VALUE, TimeUnit.MILLISECONDS).
- ❖ If producer is unable to complete all requests before the timeout expires, all unsent requests fail, and this method fails

Orderly shutdown using close

```
C StockPriceKafkaProducer.java x
StockPriceKafkaProducer main()
95 ⑨ Runtime.getRuntime().addShutdownHook(new Thread(() -> {
96
97     executorService.shutdown();
98     try {
99         executorService.awaitTermination( timeout: 200, TimeUnit.MILLISECONDS);
100        logger.info("Shutting down executorService for workers nicely");
101    } catch (InterruptedException e) {
102        logger.warn("shutting down", e);
103    }
104
105    logger.info("Flushing producer");
106    producer.flush();
107    logger.info("Closing producer");
108    producer.close();
109
110    if (!executorService.isShutdown()) {
111        logger.info("Forcing shutdown of workers");
112        executorService.shutdownNow();
113    }
114});
```

Wait for clean close

```
c StockPriceKafkaProducer.java x
StockPriceKafkaProducer main()
95  Runtime.getRuntime().addShutdownHook(new Thread(() -> {
96
97      executorService.shutdown();
98
99      try {
100          executorService.awaitTermination( timeout: 200, TimeUnit.MILLISECONDS );
101          logger.info("Flushing and closing producer");
102          producer.flush();
103          producer.close( timeout: 10_000, TimeUnit.MILLISECONDS );
104      } catch (InterruptedException e) {
105          logger.warn("shutting down", e);
106      }
107  }));

```

KafkaProducer partitionsFor() method

- ❖ partitionsFor(topic) returns meta data for partitions
- ❖ **public** List<PartitionInfo> partitionsFor(String topic)
- ❖ Get partition metadata for give topic
- ❖ Producers that do their own partitioning would use this
 - ❖ for custom partitioning
- ❖ PartitionInfo(String topic, int partition, Node leader, Node[] replicas, Node[] inSyncReplicas)
 - ❖ Node(int id, String host, int port, optional String rack)

StockPrice Producer Java Example

StockPrice App to demo Advanced Producer

- ❖ **StockPrice** - holds a stock price has a name, dollar, and cents
- ❖ **StockPriceKafkaProducer** - Configures and creates **KafkaProducer<String, StockPrice>**, **StockSender** list, ThreadPool (ExecutorService), starts **StockSender** runnable into thread pool
- ❖ **StockAppConstants** - holds topic and broker list
- ❖ **StockPriceSerializer** - can serialize a **StockPrice** into **byte[]**
- ❖ **StockSender** - generates somewhat random stock prices for a given **StockPrice** name, Runnable, 1 thread per StockSender
 - ❖ Shows using **KafkaProducer** from many threads

StockPrice domain object

```
c StockPrice.java x
StockPrice
1 package com.cloudurable.kafka.producer.model;
2
3 import io.advantageous.boon.json.JsonFactory;
4
5 public class StockPrice {
6
7     private final int dollars;
8     private final int cents;
9     private final String name;
10
11    public String toJson() {
12        return "{" +
13            "\"dollars\": " + dollars +
14            ", \"cents\": " + cents +
15            ", \"name\": \"\" + name + '\"' +
16            '}';
17    }
18}
```

- ❖ has name
- ❖ dollars
- ❖ cents
- ❖ converts itself to JSON

StockPriceKafkaProducer

- ❖ Import classes and setup logger
- ❖ Create ***createProducer*** method to create ***KafkaProducer*** instance
- ❖ Create ***setupBootstrapAndSerializers*** to initialize bootstrap servers, client id, key serializer and custom serializer (***StockPriceSerializer***)
- ❖ Write ***main()*** method - creates ***producer***, create ***StockSender*** list passing each instance a ***producer***, creates a thread pool so every stock sender gets its own thread, runs each ***stockSender*** in its own thread

StockPriceKafkaProducer imports, createProducer

Tos

```
StockPriceKafkaProducer main()
1 package com.cloudurable.kafka.producer;
2
3 import com.cloudurable.kafka.StockAppConstants;
4 import com.cloudurable.kafka.producer.model.StockPrice;
5 import io.advantageous.boon.core.Lists;
6 import org.apache.kafka.clients.producer.*;
7 import org.apache.kafka.common.serialization.StringSerializer;
8 import org.slf4j.Logger;
9 import org.slf4j.LoggerFactory;                                Import classes and
                                                               setup logger
10
11 import java.util.List;                                     createProducer used to
12 import java.util.Properties;                               create a KafkaProducer
13 import java.util.concurrent.ExecutorService;
14 import java.util.concurrent.Executors;
15 import java.util.concurrent.TimeUnit;                      createProducer() calls
                                                               setupBootstrapAnd Serializers()
16
17 public class StockPriceKafkaProducer {                     createProducer()
18     private static final Logger logger =                  setupBootstrapAnd Serializers()
19         LoggerFactory.getLogger(StockPriceKafkaProducer.class);
20
21     private static Producer<String, StockPrice> createProducer() {
22         final Properties props = new Properties();
23         setupBootstrapAndSerializers(props);
24         return new KafkaProducer<>(props);
25     }
}
```

Configure Producer Bootstrap and Serializer

```
StockPriceKafkaProducer
27     private static void setupBootstrapAndSerializers(Properties props) {
28         props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
29                   StockAppConstants.BOOTSTRAP_SERVERS);
30         props.put(ProducerConfig.CLIENT_ID_CONFIG, "StockPriceKafkaProducer");
31         props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
32                   StringSerializer.class.getName());
33
34         //Custom Serializer - config "value.serializer"
35         props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
36                   StockPriceSerializer.class.getName());
37     }
```

- ❖ Create ***setupBootstrapAndSerializers*** to initialize bootstrap servers, client id, key serializer and custom serializer (***StockPriceSerializer***)
- ❖ ***StockPriceSerializer*** will serialize ***StockPrice*** into bytes

StockPriceKafkaProducer.main()

```
StockPriceKafkaProducer.java x
StockPriceKafkaProducer main()
39 >     public static void main(String... args)
40             throws Exception {
41
42     //Create Kafka Producer
43     final Producer<String, StockPrice> producer = createProducer();
44
45     //Create StockSender list
46     final List<StockSender> stockSenders = getStockSenderList(producer);
47
48     //Create a thread pool so every stock sender gets its own.
49     final ExecutorService executorService =
50         Executors.newFixedThreadPool(stockSenders.size());
51
52     //Run each stock sender in its own thread.
53     stockSenders.forEach(executorService::submit);
54
55 }
```

- ❖ *main* method - creates *producer*,
- ❖ create *StockSender* list passing each instance a *producer*
- ❖ creates a thread pool (executorService)
- ❖ every StockSender runs in its own thread

StockAppConstants

```
1 package com.cloudurable.kafka;  
2  
3 public class StockAppConstants {  
4     public final static String TOPIC = "stock-prices";  
5     public final static String BOOTSTRAP_SERVERS =  
6         "localhost:9092,localhost:9093,localhost:9094";  
7  
8 }
```

- ❖ topic name for Producer example
- ❖ List of bootstrap servers

StockPriceKafkaProducer.getStockSenderList

Tos

```
StockPriceKafkaProducer getStockSenderList()
57 @  private static List<StockSender> getStockSenderList(
58     final Producer<String, StockPrice> producer) {
59     return Lists.list(
60         new StockSender(StockAppConstants.TOPIC,
61             new StockPrice( name: "IBM", dollars: 100, cents: 99),
62             new StockPrice( name: "IBM", dollars: 50, cents: 10),
63             producer,
64             delayMinMs: 1, delayMaxMs: 10
65         ),
66         new StockSender(
67             StockAppConstants.TOPIC,
68             new StockPrice( name: "SUN", dollars: 100, cents: 99),
69             new StockPrice( name: "SUN", dollars: 50, cents: 10),
70             producer,
71             delayMinMs: 1, delayMaxMs: 10
72         ),
73         new StockSender(
74             StockAppConstants.TOPIC,
75             new StockPrice( name: "GOOG", dollars: 500, cents: 99),
76             new StockPrice( name: "GOOG", dollars: 400, cents: 10),
77             producer,
78             delayMinMs: 1, delayMaxMs: 10
79         ),
80         new StockSender(
81             StockAppConstants.TOPIC,
82             new StockPrice( name: "INEL", dollars: 100, cents: 99),
83             new StockPrice( name: "INEL", dollars: 50, cents: 10),
84             producer,
85             delayMinMs: 1, delayMaxMs: 10
)
```

StockPriceSerializer

```
1 package com.cloudurable.kafka.producer;
2 import com.cloudurable.kafka.producer.model.StockPrice;
3 import org.apache.kafka.common.serialization.Serializer;
4 import java.nio.charset.StandardCharsets;
5 import java.util.Map;
6
7 public class StockPriceSerializer implements Serializer<StockPrice> {
8
9     @Override
10    public byte[] serialize(String topic, StockPrice data) {
11        return data.toJson().getBytes(StandardCharsets.UTF_8);
12    }
13
14    @Override
15    public void configure(Map<String, ?> configs, boolean isKey) {
16    }
17
18    @Override
19    public void close() {
20    }
21}
```

❖ Converts *StockPrice* into byte array

StockSender run()

```
StockSender run()
38
39 ①↑ public void run() {
40      final Random random = new Random(System.currentTimeMillis());
41      int sentCount = 0;
42
43      while (true) {
44          sentCount++;
45          final ProducerRecord <String, StockPrice> record =
46              createRandomRecord(random);
47          final int delay = randomIntBetween(random, delayMaxMs, delayMinMs);
48
49          try {
50              final Future<RecordMetadata> future = producer.send(record);
51              if (sentCount % 100 == 0) {displayRecordMetaData(record, future);}
52              Thread.sleep(delay);
53          } catch (InterruptedException e) {
54              if (Thread.interrupted()) {
55                  break;
56              }
57          } catch (ExecutionException e) {
58              logger.error("problem sending record to producer", e);
59          }
60      }
61 }
```

- ❖ In loop, creates random record, send record, waits random time

StockSender displayRecordMetaData

```
StockSender displayRecordMetaData()
62
63     private void displayRecordMetaData(final ProducerRecord<String, StockPrice> record,
64                                         final Future<RecordMetadata> future)
65                                         throws InterruptedException, ExecutionException {
66         final RecordMetadata recordMetadata = future.get();
67         logger.info(String.format("\n\t\tkey=%s, value=%s " +
68                               "\n\t\tsent to topic=%s part=%d off=%d at time=%s",
69                     record.key(),
70                     record.value().toJson(),
71                     recordMetadata.topic(),
72                     recordMetadata.partition(),
73                     recordMetadata.offset(),
74                     new Date(recordMetadata.timestamp())))
75     }
76 }
```

- ❖ Every 100 records `displayRecordMetaData` gets called
- ❖ Prints out record info, and recordMetadata info:
 - ❖ key, JSON value, topic, partition, offset, time
 - ❖ uses Future from call to `producer.send()`

Run All 3 Brokers

```
1 #!/usr/bin/env bash
2 CONFIG=`pwd`/config
3 cd ~/kafka-training
4 ## Run Kafka
5 kafka/bin/kafka-server-start.sh \
6     "$CONFIG/server-0.properties"
7
```

```
start-2nd-server.sh x
1 #!/usr/bin/env bash
2 CONFIG=`pwd`/config
3 cd ~/kafka-training
4 ## Run Kafka
5 kafka/bin/kafka-server-start.sh \
6     "$CONFIG/server-1.properties"
7
```

```
start-3rd-server.sh x
1 #!/usr/bin/env bash
2 CONFIG=`pwd`/config
3 cd ~/kafka-training
4 ## Run Kafka
5 kafka/bin/kafka-server-start.sh \
6     "$CONFIG/server-2.properties"
7
```



```
server-1.properties x
1 broker.id=1
2 port=9093
3 log.dirs=./logs/kafka-1
4 min.insync.replicas=3
5 compression.type=producer
```

```
server-2.properties x
1 broker.id=2
2 port=9094
3 log.dirs=./logs/kafka-2
4 min.insync.replicas=3
5 compression.type=producer
6 auto.create.topics.enable=false
7 message.max.bytes=65536
8 replica.lag.time.max.ms=5000
9 delete.topic.enable=true
```

Run create-topic.sh script

Tos

```
create-topic.sh x
1 #!/usr/bin/env bash
2
3 cd ~/kafka-training
4
5 kafka/bin/kafka-topics.sh \
6   --create \
7   --zookeeper localhost:2181 \
8   --replication-factor 3 \
9   --partitions 3 \
10  --topic stock-prices
11
```

Name of the topic is stock-prices

Three partitions

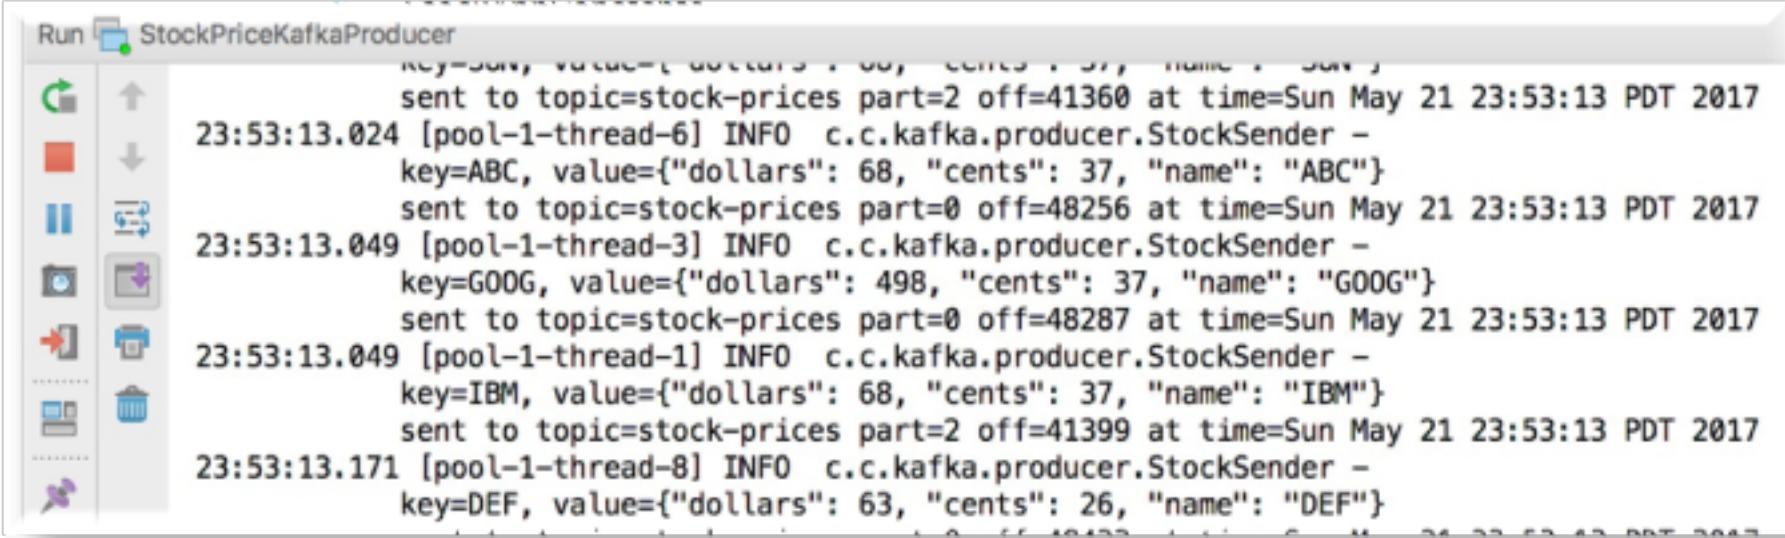
Replication factor of three

Terminal

+ Local Local (1) Local (2) Local (3) Local (4)

✗ \$./bin/create-topic.sh
Created topic "stock-prices".

Run StockPriceKafkaProducer



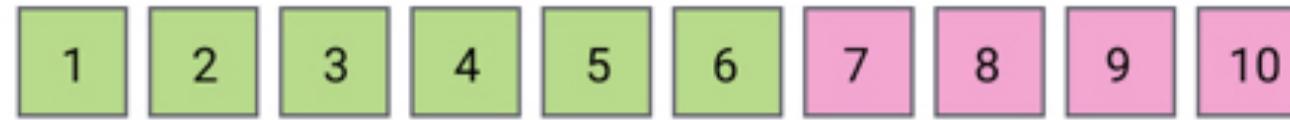
The screenshot shows a run window titled "Run StockPriceKafkaProducer". The log output displays several INFO messages from the `c.c.kafka.producer.StockSender` class. Each message indicates a stock price being sent to the "stock-prices" topic. The log entries are as follows:

```
key=ABC, value={"dollars": 68, "cents": 37, "name": "ABC"}  
sent to topic=stock-prices part=2 off=41360 at time=Sun May 21 23:53:13 PDT 2017  
23:53:13.024 [pool-1-thread-6] INFO c.c.kafka.producer.StockSender -  
key=GOOG, value={"dollars": 498, "cents": 37, "name": "GOOG"}  
sent to topic=stock-prices part=0 off=48256 at time=Sun May 21 23:53:13 PDT 2017  
23:53:13.049 [pool-1-thread-3] INFO c.c.kafka.producer.StockSender -  
key=IBM, value={"dollars": 68, "cents": 37, "name": "IBM"}  
sent to topic=stock-prices part=0 off=48287 at time=Sun May 21 23:53:13 PDT 2017  
23:53:13.049 [pool-1-thread-1] INFO c.c.kafka.producer.StockSender -  
key=DEF, value={"dollars": 63, "cents": 26, "name": "DEF"}  
sent to topic=stock-prices part=2 off=41399 at time=Sun May 21 23:53:13 PDT 2017  
23:53:13.171 [pool-1-thread-8] INFO c.c.kafka.producer.StockSender -
```

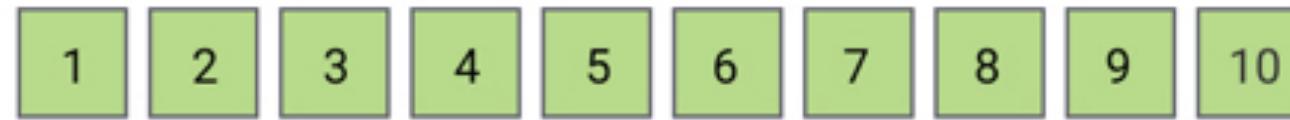
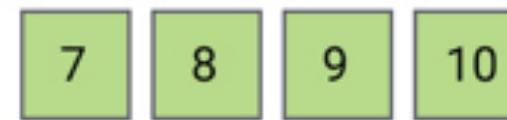
- ❖ Run **StockPriceKafkaProducer** from the IDE
- ❖ You should see log messages from **StockSender(s)** with **StockPrice** name, JSON value, partition, offset, and time

idempotent producer

Producer Resends on Failure



enable.idempotence=true.



Limitation 1: Ack=All

Limitation 2: max.in.flight.requests.per.connection <= 5

Consumer reads duplicates



How it Works

- M1 (PID: 1, SN: 1) - written to partition. For PID 1, Max SN=1
- M2 (PID: 1, SN: 2) - written to partition. For PID 1, Max SN=2
- M3 (PID: 1, SN: 3) - written to partition. For PID 1, Max SN=3
- M4 (PID: 1, SN: 4) - written to partition. For PID 1, Max SN=4
- M5 (PID: 1, SN: 5) - written to partition. For PID 1, Max SN=5
- M6 (PID: 1, SN: 6) - written to partition. For PID 1, Max SN=6
- M4 (PID: 1, SN: 4) - rejected, SN <= Max SN
- M5 (PID: 1, SN: 5) - rejected, SN <= Max SN
- M6 (PID: 1, SN: 6) - rejected, SN <= Max SN
- M7 (PID: 1, SN: 7) - written to partition. For PID 1, Max SN=7
- M8 (PID: 1, SN: 8) - written to partition. For PID 1, Max SN=8
- M9 (PID: 1, SN: 9) - written to partition. For PID 1, Max SN=9
- M10 (PID: 1, SN: 10) - written to partition. For PID 1, Max SN=10

How it Works

Each producer gets assigned a **Producer Id (PID)** and it includes its PID every time it sends messages to a broker.

Additionally, each message gets a **monotonically increasing sequence number**.

A separate sequence is maintained for each topic partition that a producer sends messages to.

On the broker side, on a per partition basis, it keeps track of the largest PID-Sequence Number combination is has successfully written.

When a lower sequence number is received, it is discarded

transactional producer

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("transactional.id", "my-transactional-id");
Producer<String, String> producer = new KafkaProducer<>(props, new StringSerializer(), new StringSerialize

producer.initTransactions();

try {
    producer.beginTransaction();
    for (int i = 0; i < 100; i++)
        producer.send(new ProducerRecord<>"my-topic", Integer.toString(i), Integer.toString(i)));
    producer.commitTransaction();
} catch (ProducerFencedException | OutOfOrderSequenceException | AuthorizationException e) {
    // We can't recover from these exceptions, so our only option is to close the producer and exit.
    producer.close();
} catch (KafkaException e) {
    // For all other exceptions, just abort the transaction and try again.
    producer.abortTransaction();
}
producer.close();
```

Message serialization using Avro

Avro relies on schemas so as to provide efficient serialization of the data.

The schema is written in JSON format and describes the fields and their types.

There are 2 cases:

- when serializing to a file, the schema is written to the file
- in RPC - such as between Kafka and Spark - both systems should know the schema prior to exchanging data, or they could exchange the schema during the connection handshake.

What makes Avro handy is that you do not need to generate data classes.

Message serialization using Avro

Add dependency & Schema

```
<dependency>
  <groupId>org.apache.avro</groupId>
  <artifactId>avro</artifactId>
  <version>1.8.0</version>
</dependency>
```

```
{
  "fields": [
    { "name": "str1", "type": "string" },
    { "name": "str2", "type": "string" },
    { "name": "int1", "type": "int" }
  ],
  "name": "myrecord",
  "type": "record"
}
```

Message serialization using Avro - The Producer

```
public static void main(String[] args) throws InterruptedException {
    Properties props = new Properties();
    props.put("bootstrap.servers", "localhost:9092");
    props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
    props.put("value.serializer", "org.apache.kafka.common.serialization.ByteArraySerializer");

    Schema.Parser parser = new Schema.Parser();
    Schema schema = parser.parse(USER_SCHEMA);
    Injection<GenericRecord, byte[]> recordInjection = GenericAvroCodecs.toBinary(schema);

    KafkaProducer<String, byte[]> producer = new KafkaProducer<>(props);
    public static final String USER_SCHEMA = "{"
        + "\"type\":\"record\","
        + "\"name\":\"myrecord\","
        + "\"fields\":["
        + "  { \"name\":\"str1\", \"type\":\"string\" },"
        + "  { \"name\":\"str2\", \"type\":\"string\" },"
        + "  { \"name\":\"int1\", \"type\":\"int\" }"
        + "]}";

    for (int i = 0; i < 1000; i++) {
        GenericData.Record avroRecord = new GenericData.Record(schema);
        avroRecord.put("str1", "Str 1-" + i);
        avroRecord.put("str2", "Str 2-" + i);
        avroRecord.put("int1", i);

        byte[] bytes = recordInjection.apply(avroRecord);

        ProducerRecord<String, byte[]> record = new ProducerRecord<>("mytopic", bytes);
        producer.send(record);

        Thread.sleep(250);
    }
    producer.close();
}
```

Lab : Java API – Producer-II a API