

Kafka Consumers Advance

- ❖ A Consumer client
 - ❖ consumes records from Kafka cluster
- ❖ Automatically handles Kafka broker failure
 - ❖ adapts as topic partitions leadership moves in Kafka cluster
- ❖ Works with Kafka broker to form consumers groups and load balance consumers
- ❖ Consumer maintains connections to Kafka brokers in cluster
- ❖ Use ***close()*** method to not leak resources
- ❖ NOT thread-safe

Consumer: createConsumer / Consumer Config

```
SimpleStockPriceConsumer.java x
SimpleStockPriceConsumer

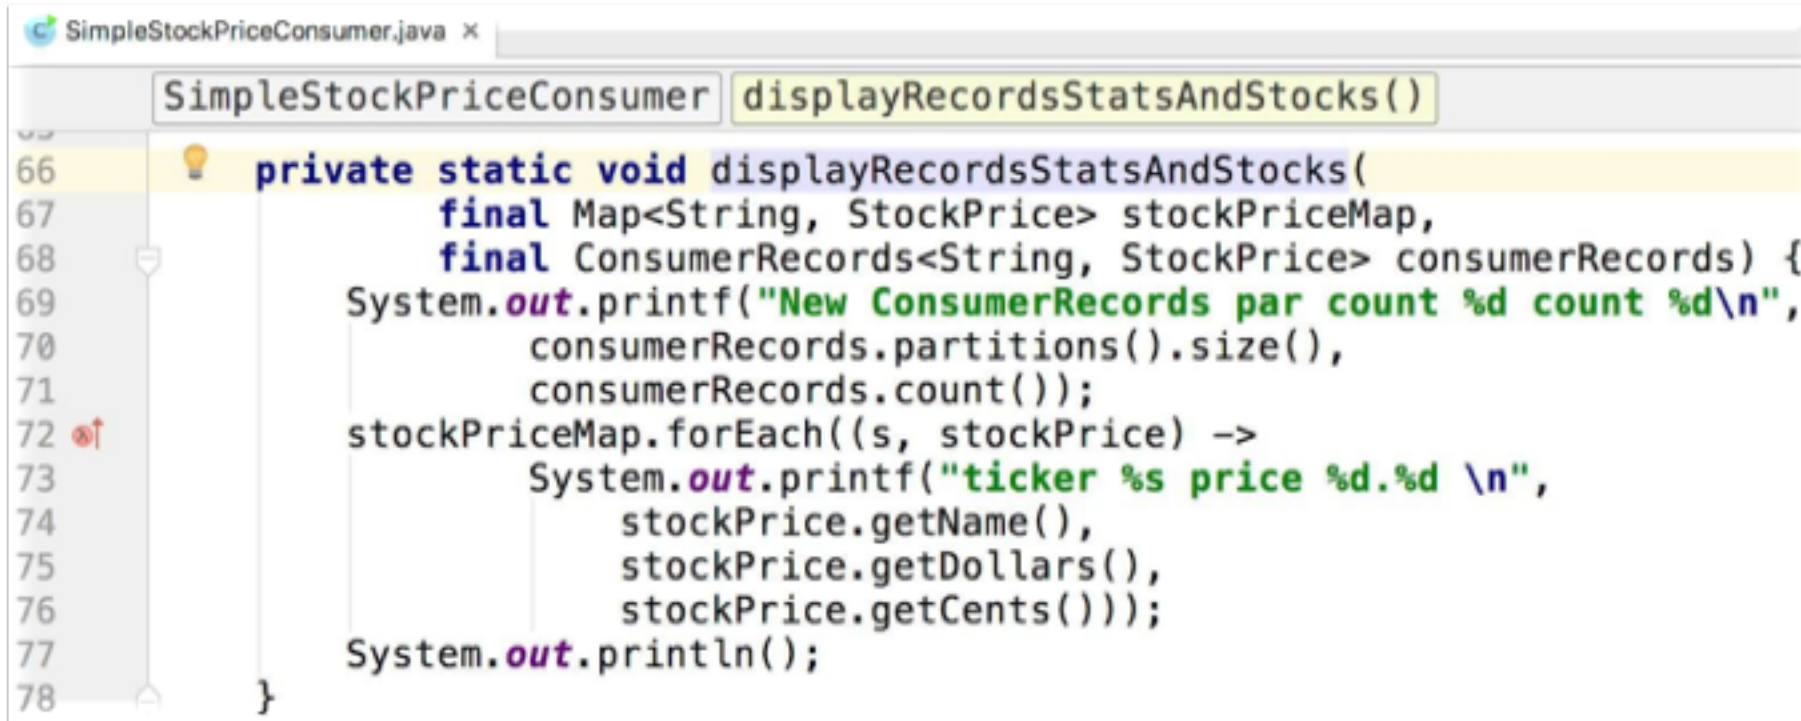
11
12 ▶ public class SimpleStockPriceConsumer {
13
14     private static Consumer<String, StockPrice> createConsumer() {
15         final Properties props = new Properties();
16         props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,
17                 StockAppConstants.BOOTSTRAP_SERVERS);
18         props.put(ConsumerConfig.GROUP_ID_CONFIG,
19                 "KafkaExampleConsumer");
20         props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
21                 StringDeserializer.class.getName());
22         //Custom Deserializer
23         props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
24                 StockDeserializer.class.getName());
25         props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, 500);
26         // Create the consumer using props.
27         final Consumer<String, StockPrice> consumer =
28             new KafkaConsumer<>(props);
29         // Subscribe to the topic.
30         consumer.subscribe(Collections.singletonList(
31             StockAppConstants.TOPIC));
32         return consumer;
33     }
```

- ❖ Similar to other Consumer examples so far
- ❖ Subscribes to **stock-prices** topic
- ❖ Has custom serializer

SimpleStockPriceConsumer.runConsumer

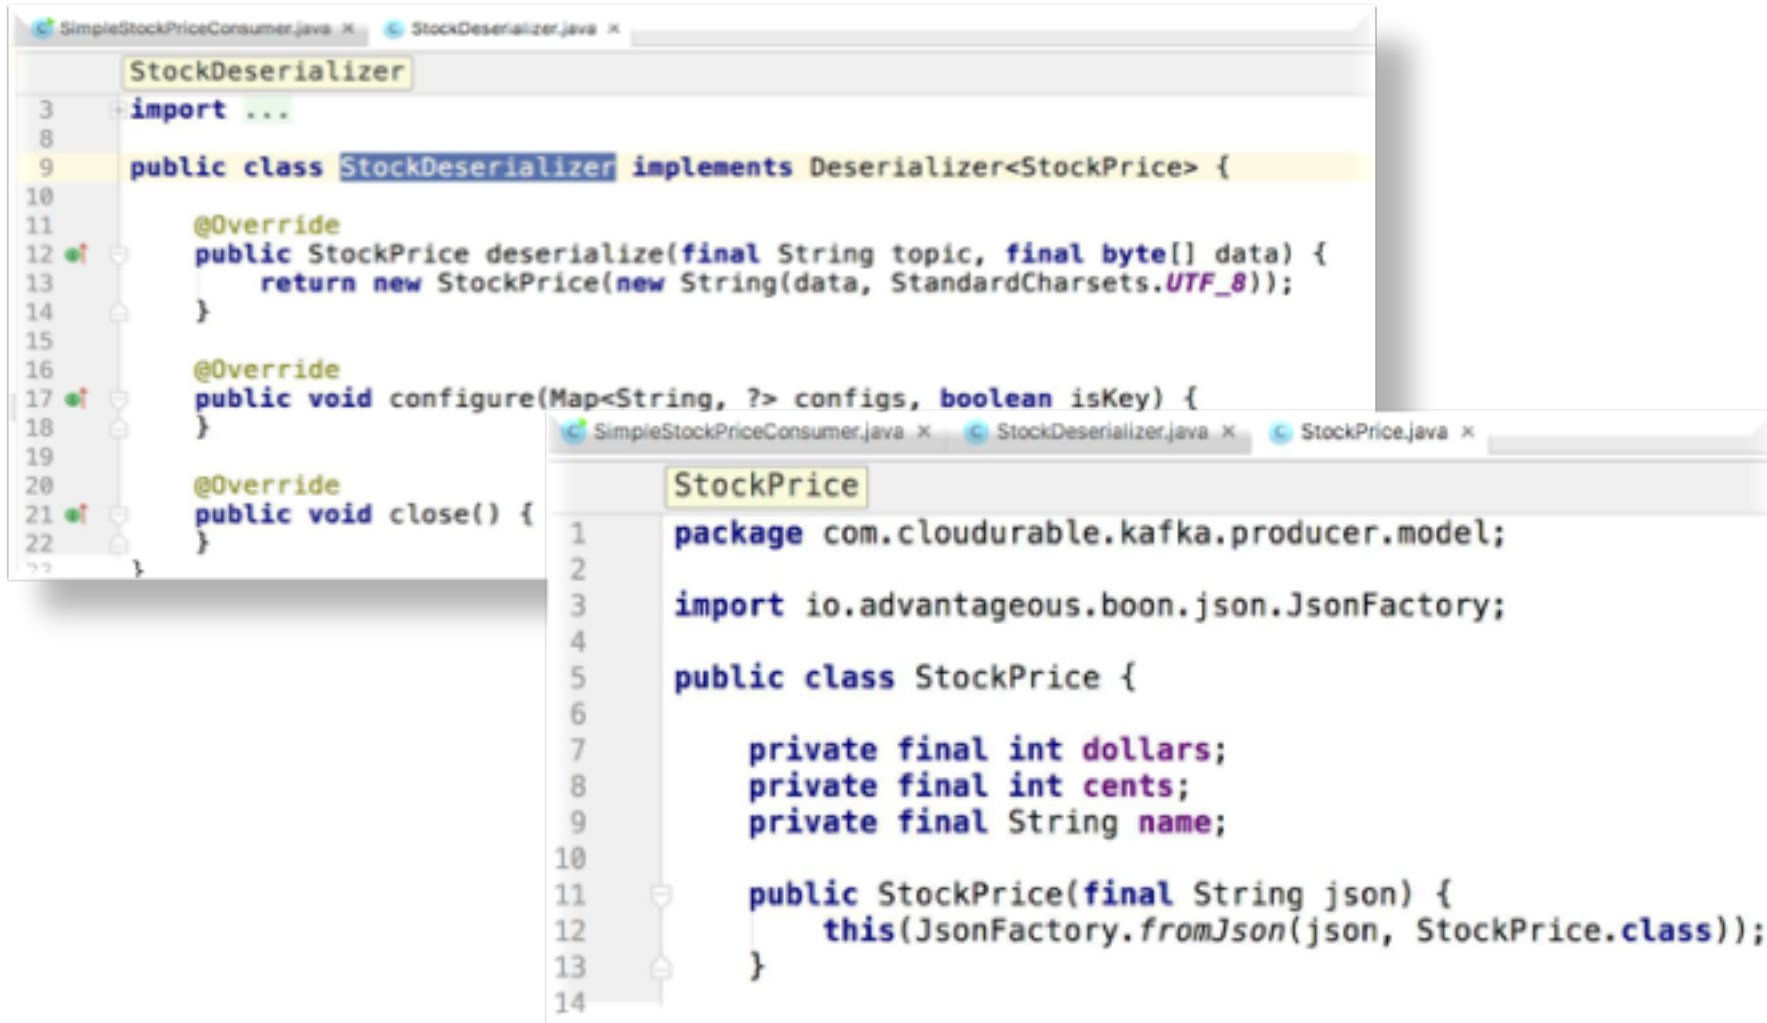
```
SimpleStockPriceConsumer.java x
SimpleStockPriceConsumer
35
36 static void runConsumer() throws InterruptedException {
37     final Consumer<String, StockPrice> consumer = createConsumer();
38     final Map<String, StockPrice> map = new HashMap<>();
39     try {
40         final int giveUp = 1000; int noRecordsCount = 0;
41         int readCount = 0;
42         while (true) {
43             final ConsumerRecords<String, StockPrice> consumerRecords =
44                 consumer.poll( timeout: 1000);
45             if (consumerRecords.count() == 0) {
46                 noRecordsCount++;
47                 if (noRecordsCount > giveUp) break;
48                 else continue;
49             }
50             readCount++;
51             consumerRecords.forEach(record -> {
52                 map.put(record.key(), record.value());
53             });
54             if (readCount % 100 == 0) {
55                 displayRecordsStatsAndStocks(map, consumerRecords);
56             }
57             consumer.commitAsync();
58         }
59     }
```

- ❖ Drains topic; Creates map of current stocks; Calls *displayRecordsStatsAndStocks()*



```
SimpleStockPriceConsumer.java x
SimpleStockPriceConsumer displayRecordsStatsAndStocks()
66 private static void displayRecordsStatsAndStocks(
67     final Map<String, StockPrice> stockPriceMap,
68     final ConsumerRecords<String, StockPrice> consumerRecords) {
69     System.out.printf("New ConsumerRecords par count %d count %d\n",
70         consumerRecords.partitions().size(),
71         consumerRecords.count());
72     stockPriceMap.forEach((s, stockPrice) ->
73         System.out.printf("ticker %s price %d.%d \n",
74             stockPrice.getName(),
75             stockPrice.getDollars(),
76             stockPrice.getCents()));
77     System.out.println();
78 }
```

- ❖ Prints out size of each partition read and total record count
- ❖ Prints out each stock at its current price



The image shows two overlapping Java code editors. The top editor displays the `StockDeserializer` class, which implements the `Deserializer<StockPrice>` interface. It includes methods for `deserialize`, `configure`, and `close`. The bottom editor displays the `StockPrice` class, which is a simple data object with fields for `dollars`, `cents`, and `name`, and a constructor that takes a JSON string.

```
SimpleStockPriceConsumer.java x StockDeserializer.java x
StockDeserializer
3 import ...
8
9 public class StockDeserializer implements Deserializer<StockPrice> {
10
11     @Override
12     public StockPrice deserialize(final String topic, final byte[] data) {
13         return new StockPrice(new String(data, StandardCharsets.UTF_8));
14     }
15
16     @Override
17     public void configure(Map<String, ?> configs, boolean isKey) {
18     }
19
20     @Override
21     public void close() {
22     }
23 }

SimpleStockPriceConsumer.java x StockDeserializer.java x StockPrice.java x
StockPrice
1 package com.cloudurable.kafka.producer.model;
2
3 import io.advantageous.boon.json.JsonFactory;
4
5 public class StockPrice {
6
7     private final int dollars;
8     private final int cents;
9     private final String name;
10
11     public StockPrice(final String json) {
12         this(JsonFactory.fromJson(json, StockPrice.class));
13     }
14 }
```

- ❖ If using partition assignment, you must handle cases where partition assignments change
- ❖ Pass ***ConsumerRebalanceListener*** instance in call to ***kafkaConsumer.subscribe(Collection, ConsumerRebalanceListener)*** and ***kafkaConsumer.subscribe(Pattern, ConsumerRebalanceListener)***.
- ❖ when partitions taken from consumer, commit its offset for partitions by implementing ***ConsumerRebalanceListener.onPartitionsRevoked(Collection)***
- ❖ When partitions are assigned to consumer, look up offset for new partitions and correctly initialize consumer to that position by implementing ***ConsumerRebalanceListener.onPartitionsAssigned(Collection)***

```
// Subscribe to the topic.  
consumer.subscribe(Collections.singletonList(  
    StockAppConstants.TOPIC),  
    new SeekToConsumerRebalanceListener(consumer, seekTo, location));
```


Controlling Consumers Position Example

```
SeekToConsumerRebalanceListener onPartitionsAssigned()
8  import java.util.Collection;
9
10 public class SeekToConsumerRebalanceListener implements ConsumerRebalanceListener {
11     private final Consumer<String, StockPrice> consumer;
12     private final SeekTo seekTo; private boolean done;
13     private final long location;
14     private final long startTime = System.currentTimeMillis();
15     public SeekToConsumerRebalanceListener(final Consumer<String, StockPrice> consumer
20
21     @Override
22     public void onPartitionsAssigned(final Collection<TopicPartition> partitions) {
23         if (done) return;
24         else if (System.currentTimeMillis() - startTime > 30_000) {
25             done = true;
26             return;
27         }
28         switch (seekTo) {
29             case END: //Seek to end
30                 consumer.seekToEnd(partitions);
31                 break;
32             case START: //Seek to start
33                 consumer.seekToBeginning(partitions);
34                 break;
35             case LOCATION: //Seek to a given location
36                 partitions.forEach(topicPartition ->
37                     consumer.seek(topicPartition, location));
38                 break;
39         }
40     }
```


Controlling Consumers Position Example

```
private static KafkaConsumer<String, String> startConsumer(String name) {
    Properties consumerProps = ExampleConfig.getConsumerProps();
    KafkaConsumer<String, String> consumer = new KafkaConsumer<>(consumerProps);
    consumer.subscribe(Collections.singleton("example-topic-2020-6-24"),
        new ConsumerRebalanceListener() {
            @Override
            public void onPartitionsRevoked(Collection<TopicPartition> partitions) {
                System.out.printf("onPartitionsRevoked - consumerName: %s, partitions: %s%n", name,
                    formatPartitions(partitions));
            }

            @Override
            public void onPartitionsAssigned(Collection<TopicPartition> partitions) {
                System.out.printf("onPartitionsAssigned - consumerName: %s, partitions: %s%n", name,
                    formatPartitions(partitions));
            }
        });
    System.out.printf("starting consumerName: %s%n", name);
    consumer.poll(Duration.ofSeconds(10));
    System.out.printf("closing consumerName: %s%n", name);
    consumer.close();
    return consumer;
}
```

- ❖ Consumers join consumer group after subscribe and then ***poll()*** is called
- ❖ Automatically, consumer sends periodic heartbeats to Kafka brokers server
- ❖ If consumer crashes or unable to send heartbeats for a duration of ***session.timeout.ms***, then consumer is deemed dead and its partitions are reassigned

- ❖ Instead of subscribing to the topic using `subscribe`, you can call ***assign(Collection)*** with the full topic partition list

```
String topic = "log-replication";
```

```
TopicPartition part0 = new TopicPartition(topic, 0);
```

```
TopicPartition part1 = new TopicPartition(topic, 1);
```

```
consumer.assign(Arrays.asList(part0, part1));
```

- ❖ Using consumer as before with ***poll()***
- ❖ Manual partition assignment negates use of group coordination, and auto consumer fail over
 - Each consumer acts independently even if in a consumer group (use unique group id to avoid confusion)
- ❖ You have to use ***assign()*** or ***subscribe()*** but not both

- ❖ Calling ***poll()*** marks consumer as alive
 - ❖ If consumer continues to call ***poll()***, then consumer is alive and in consumer group and gets messages for partitions assigned (has to call before every ***max.poll.interval.ms interval***)
 - ❖ Not calling ***poll()***, even if consumer is sending heartbeats, consumer is still considered dead
- ❖ Processing of records from ***poll*** has to be faster than ***max.poll.interval.ms*** interval or your consumer could be marked dead!
- ❖ ***max.poll.records*** is used to limit total records returned from a ***poll*** call - easier to predict max time to process records on each ***poll*** interval

- ❖ ***At most once***
 - ❖ Messages may be lost but are never redelivered
- ❖ ***At least once***
 - ❖ Messages are never lost but may be redelivered
- ❖ ***Exactly once***
 - ❖ this is what people actually want, each message is delivered once and only once

“At-Least-Once” - Delivery Semantics

SimpleStockPriceConsumer pollRecordsAndProcess()

```
70
77     final ConsumerRecords<String, StockPrice> consumerRecords =
78         consumer.poll( timeout: 1000);
79
80     try {
81         startTransaction();           //Start DB Transaction
82
83                                     //Process the records
84         processRecords(map, consumerRecords);
85
86                                     //Commit the Kafka offset
87         consumer.commitSync();
88
89         commitTransaction();          //Commit DB Transaction
90     } catch (CommitFailedException ex) {
91         logger.error("Failed to commit sync to log", ex);
92         rollbackTransaction();        //Rollback Transaction
93     } catch (DatabaseException dte) {
94         logger.error("Failed to write to DB", dte);
95         rollbackTransaction();        //Rollback Transaction
96     }
```

“At-Most-Once” - Delivery Semantics

SimpleStockPriceConsumer pollRecordsAndProcess()

```
77      final ConsumerRecords<String, StockPrice> consumerRecords =  
78          consumer.poll( timeout: 1000);  
79  
80      try {  
81          startTransaction();           //Start DB Transaction  
82  
83          //Commit the Kafka offset  
84          consumer.commitSync();  
85  
86          //Process the records  
87          processRecords(map, consumerRecords);  
88  
89          commitTransaction();          //Commit DB Transaction  
90      } catch (CommitFailedException ex) {  
91          logger.error("Failed to commit sync to log", ex);  
92          rollbackTransaction();        //Rollback Transaction  
93      } catch (DatabaseException dte) {  
94          logger.error("Failed to write to DB", dte);  
95          rollbackTransaction();        //Rollback Transaction  
96      }
```


- ❖ Consumer do not have to use Kafka's built-in offset storage
- ❖ Consumers can choose to store offsets with processed record output to make it “exactly once” message consumption
- ❖ If Consumer output of record consumption is stored in RDBMS then storing offset in database allows committing both process record output and location (partition/offset of record) in a single transaction implementing “exactly once” messaging.
- ❖ Typically to achieve “exactly once” you store record location with output of record

Saving Topic, Offset, Partition in DB

```
DatabaseUtilities saveStockPrice()  
22  
23 public static void saveStockPrice(final StockPriceRecord stockRecord,  
24                                     final Connection connection) throws SQLException {  
25  
26     final PreparedStatement preparedStatement = getUpsertPreparedStatement(  
27         stockRecord.getName(), connection);  
28  
29  
30  
31     //Save partition, offset and topic in database.  
32     preparedStatement.setLong( parameterIndex: 1, stockRecord.getOffset());  
33     preparedStatement.setLong( parameterIndex: 2, stockRecord.getPartition());  
34     preparedStatement.setString( parameterIndex: 3, stockRecord.getTopic());  
35  
36     //Save stock price, name, dollars, and cents into database.  
37     preparedStatement.setInt( parameterIndex: 4, stockRecord.getDollars());  
38     preparedStatement.setInt( parameterIndex: 5, stockRecord.getCents());  
39     preparedStatement.setString( parameterIndex: 6, stockRecord.getName());  
40  
41     //Save the record with offset, partition, and topic.  
42     preparedStatement.execute();  
43  
44 }
```

- ❖ If implementing “*exactly once*” message semantics, then you have to manage offset positioning
 - ❖ Pass ***ConsumerRebalanceListener*** instance in call to ***kafkaConsumer.subscribe(Collection, ConsumerRebalanceListener)*** and ***kafkaConsumer.subscribe(Pattern, ConsumerRebalanceListener)***.
 - ❖ when partitions taken from consumer, commit its offset for partitions by implementing ***ConsumerRebalanceListener.onPartitionsRevoked(Collection)***
 - ❖ When partitions are assigned to consumer, look up offset for new partitions and correctly initialize consumer to that position by implementing ***ConsumerRebalanceListener.onPartitionsAssigned(Collection)***

Transaction

In order for this to work, consumers reading from transactional partitions should be configured to only read committed data.

This can be achieved by by setting the ***isolation.level=read_committed*** in the consumer's configuration.

- ❖ You can control consumption of topics using by using ***consumer.pause***(Collection) and ***consumer.resume***(Collection)
 - ❖ This pauses or resumes consumption on specified assigned partitions for future ***consumer.poll***(long) calls
- ❖ Use cases where consumers may want to first focus on fetching from some subset of assigned partitions at full speed, and only start fetching other partitions when these partitions have few or no data to consume
 - ❖ Priority queue like behavior from traditional MOM
- ❖ Other cases is stream processing if performing a join and one topic stream is getting behind another.

Message Deserialization - Avro

```
public class AvroConsumerExample {  
  
    public static void main(String[] str) throws InterruptedException {  
        System.out.println("Starting AutoOffsetAvroConsumerExample ...");  
        readMessages();  
    }  
  
    private static void readMessages() throws InterruptedException {  
        KafkaConsumer<String, byte[]> consumer = createConsumer();  
        // Assign to specific topic and partition.  
        consumer.assign(Arrays.asList(new TopicPartition("avro-topic", 0)));  
        processRecords(consumer);  
    }  
}
```

Message Deserialization - Avro

```
private static void processRecords(KafkaConsumer<String, byte[]> consumer) throws {  
while (true) {  
    ConsumerRecords<String, byte[]> records = consumer.poll(100);  
    long lastOffset = 0;  
for (ConsumerRecord<String, byte[]> record : records) {  
        GenericRecord genericRecord = AvroSupport.byteArrayToData(AvroSupport.getSchema(), record.value());  
        String firstName = AvroSupport.getValue(genericRecord, "firstName", String.class);  
        System.out.printf("\n\noffset = %d, key = %s, value = %s", record.offset(), record.key(),  
        firstName);  
        lastOffset = record.offset();  
    }  
  
    System.out.println("lastOffset read: " + lastOffset);  
    consumer.commitSync();  
}  
}
```


Message Deserialization - Avro

```
private static KafkaConsumer<String, byte[]> createConsumer() {  
    Properties props = new Properties();  
    props.put("bootstrap.servers", "localhost:9092");  
    String consumeGroup = "cg1";  
    props.put("group.id", consumeGroup);  
    props.put("enable.auto.commit", "true");  
    props.put("auto.offset.reset", "earliest");  
    props.put("auto.commit.interval.ms", "100");  
    props.put("heartbeat.interval.ms", "3000");  
    props.put("session.timeout.ms", "30000");  
    props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");  
    props.put("value.deserializer", "org.apache.kafka.common.serialization.ByteArrayDeserializer");  
    return new KafkaConsumer<String, byte[]>(props);  
}  
}
```

```
String jaasTemplate = "org.apache.kafka.common.security.scram.ScramLoginModule required username=\"%s\" password=\"%s\";";
String jaasCfg = String.format(jaasTemplate, username, password);

String serializer = StringSerializer.class.getName();
String deserializer = StringDeserializer.class.getName();
    props = new Properties();
    props.put("bootstrap.servers", brokers);
    props.put("group.id", username + "-consumer");
    props.put("enable.auto.commit", "true");
    props.put("auto.commit.interval.ms", "1000");
    props.put("auto.offset.reset", "earliest");
    props.put("session.timeout.ms", "30000");
    props.put("key.deserializer", deserializer);
    props.put("value.deserializer", deserializer);
    props.put("key.serializer", serializer);
    props.put("value.serializer", serializer);
    props.put("security.protocol", "SASL_SSL");
    props.put("sasl.mechanism", "SCRAM-SHA-256");
    props.put("sasl.jaas.config", jaasCfg);
}

public void consume() {
    KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);
    consumer.subscribe(Arrays.asList(topic));
```

Lab : Java API – Consumer – II